

## L2: Skip Lists

- Think TreeMap in Java
- Maps Key  $\rightarrow$  Value. The keys are ordered. Support
  - add
  - remove
  - update
  - lookup
- Goal: for most operations, take  $O(\log n)$  time or faster.

### Applications

- Table Mapping / DB indexes. Make a collection searchable quickly.
- Best-first Heuristics.
  - ① Pack items into bins (max cap: C). A decent heuristic is to find the best-fit bin (i.e., least remaining cap but enough for the item). Use ordered map to locate the best bin, reduce size, then put it back to the pool.
  - ② malloc: find the best free block to use.

The Theme: approximating a perfectly balanced struct.   
 (i) (ii)   
  $\left\{ \begin{array}{l} \text{in prob.} \\ \text{deterministically} \\ \text{on average,} \end{array} \right.$

Q: What comes to your mind when you want an ordered map?

- \* Balanced BST
- \* B-Tree
- \* Linked List

} not easy to implement

any good thing?

the smallest & the small elts are quick to access  
same for the largest elts

by contrast: the whole depth for trees.

Today's goal: Supercharging the linked list D/S.

## Lecture Notes on Skip Lists


Lecture 12 — March 18, 2004

Erik Demaine

- ~~Balanced tree structures we know at this point: B-trees, red-black trees, treaps.~~  
~~balanced trees~~
- Could you implement them right now? Probably, with time... but without looking up any details in a book?
- Skip lists are a simple randomized structure you'll never forget.

### Starting from scratch

- Initial goal: *just searches* — ignore updates (Insert/Delete) for now
- Simplest data structure: linked list
- Sorted linked list:  $\Theta(n)$  time
- 2 sorted linked lists:
  - Each element can appear in 1 or both lists
  - How to speed up search?
  - **Idea:** Express and local subway lines
  - **Example:** 14, 23, 34, 42, 50, 59, 66, 72, 79, 86, 96, 103, 110, 116, 125  
 (What is this sequence?)
  - Boxed values are “express” stops; others are normal stops
  - Can quickly jump from express stop to next express stop, or from any stop to next normal stop
  - Represented as two linked lists, one for express stops and one for all stops:



- Every element is in linked list 2 (LL2); some elements also in linked list 1 (LL1)
- Link equal elements between the two levels
- To search, first search in LL1 until about to go too far, then go down and search in LL2

– Cost:

$$\text{len}(\text{LL1}) + \frac{\text{len}(\text{LL2})}{\text{len}(\text{LL1})} = \text{len}(\text{LL1}) + \frac{n}{\text{len}(\text{LL1})}$$

– Minimized when


$$\text{len}(\text{LL1}) = \frac{n}{\text{len}(\text{LL1})}$$

$$\Rightarrow \text{len}(\text{LL1})^2 = n$$

$$\Rightarrow \text{len}(\text{LL1}) = \sqrt{n}$$

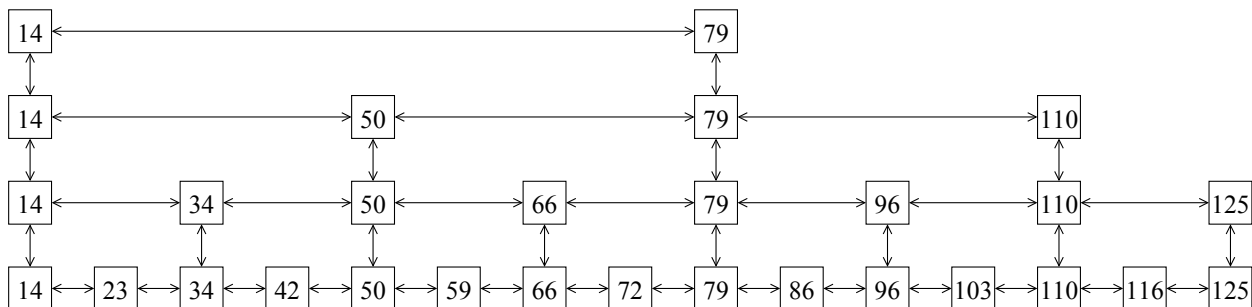
$$\Rightarrow \text{search cost} = 2\sqrt{n}$$

– Resulting 2-level structure:



- 3 linked lists:  $3 \cdot \sqrt[3]{n}$
- $k$  linked lists:  $k \cdot \sqrt[k]{n}$
- $\lg n$  linked lists:  $\lg n \cdot \sqrt[\lg n]{n} = \lg n \cdot \underbrace{n^{1/\lg n}}_{=2} = \Theta(\lg n)$   
*✓ perfect.*

– Becomes like a binary tree:



– **Example:** Search for 72

- \* Level 1: 14 too small, 79 too big; go down 14
- \* Level 2: 14 too small, 50 too small, 79 too big; go down 50
- \* Level 3: 50 too small, 66 too small, 79 too big; go down 66
- \* Level 4: 66 too small, 72 spot on

Maintaining this exactly is too rigid. Relax by flipping coins.

## Insert


- New element should certainly be added to bottommost level  
(Invariant: Bottommost list contains all elements)
- Which other lists should it be added to?  
(Is this the entire balance issue all over again?)
- **Idea:** Flip a coin
  - With what probability should it go to the next level?
  - To mimic a balanced binary tree, we'd like half of the elements to advance to the next-to-bottommost level
  - So, when you insert an element, flip a fair coin
  - If heads: add element to next level up, and flip another coin (repeat)
- Thus, on average:
  - 1/2 the elements go up 1 level
  - 1/4 the elements go up 2 levels
  - 1/8 the elements go up 3 levels
  - Etc.
- Thus, “approximately even”

## Example

- Get out a real coin and try an example
- You should put a special value  $-\infty$  at the beginning of each list, and always promote this special value to the highest level of promotion
- This forces the leftmost element to be present in every list, which is necessary for searching

... many coins are flipped ...  
(Isn't this easy?)

- The result is a skip list.
- It probably isn't as balanced as the ideal configurations drawn above.
- It's clearly good on average.
- Claim it's really really good, almost always.



## Analysis: Claim of With High Probability

- **Theorem:** *With high probability, every search costs  $\Theta(\lg n)$  in a skip list with  $n$  elements*
- What do we need to do to prove this? [Calculate the probability, and show that it's high!]
- We need to define the notion of “with high probability”; this is a powerful technical notion, used throughout randomized algorithms
- **Informal definition:** An event occurs *with high probability* if, for any  $\alpha \geq 1$ , there is an appropriate choice of constants for which  $E$  occurs with probability at least  $1 - O(1/n^\alpha)$
- In reality, the constant hidden within  $\Theta(\lg n)$  in the theorem statement actually depends on  $\alpha$   $\alpha$
- **Precise definition:** A (parameterized) event  $E_\alpha$  occurs *with high probability* if, for any  $\alpha \geq 1$ ,  $E_\alpha$  occurs with probability at least  $1 - c_\alpha/n^\alpha$ , where  $c_\alpha$  is a “constant” depending only on  $\alpha$ .
- The term  $O(1/n^\alpha)$  or more precisely  $c_\alpha/n^\alpha$  is called the *error probability*
- The idea is that the error probability can be made very very very small by setting  $\alpha$  to something big, e.g., 100

## Analysis: Warmup

- **Lemma:** With high probability, skip list with  $n$  elements has  $O(\lg n)$  levels
- (In fact, the number of levels is  $\Theta(\log n)$ , but we only need an upper bound.)
- **Proof:**

- $\Pr[\text{element } x \text{ is in more than } c \lg n \text{ levels}] = 1/2^{c \lg n} = 1/n^c$
- Recall Boole's inequality / union bound:

$$\Pr[E_1 \cup E_2 \cup \dots \cup E_n] \leq \Pr[E_1] + \Pr[E_2] + \dots + \Pr[E_n]$$

- Applying this inequality:  
 $\Pr[\text{any element is in more than } c \lg n \text{ levels}] \leq n \cdot 1/n^c = 1/n^{c-1}$
- Thus, error probability is polynomially small and exponent ( $\alpha = c - 1$ ) can be made arbitrarily large by appropriate choice of constant in level bound of  $O(\lg n)$

## Analysis: Proof of Theorem

- **Cool idea:** Analyze search backwards—from leaf to root
  - Search starts at leaf (element in bottommost level)
  - At each node visited:
    - \* If node wasn't promoted higher (got TAILS here), then we go [came from] left
    - \* If node wasn't promoted higher (got HEADS here), then we go [came from] top
  - Search stops at root of tree
- Know height is  $O(\lg n)$  with high probability; say it's  $c \lg n$
- Thus, the number of “up” moves is at most  $c \lg n$  with high probability
- Thus, search cost is at most the following quantity:
 

How many times do we need to flip a coin to get  $c \lg n$  heads?
- Intuitively,  $\Theta(\lg n)$

## Analysis: Coin Flipping

$10c \cdot \lg n$

- **Claim:** Number of flips till  $c \lg n$  heads is  $\Theta(\lg n)$  with high probability
- Again, constant in  $\Theta(\lg n)$  bound will depend on  $\alpha$
- **Proof of claim:**

– Say we make  $10c \lg n$  flips

– When are there at least  $c \lg n$  heads?


$$\text{Pr}[\text{exactly } c \lg n \text{ heads}] = \underbrace{\binom{10c \lg n}{c \lg n}}_{\substack{\text{orders} \\ \text{HHHTTT vs. HTHTHT}}} \cdot \underbrace{\left(\frac{1}{2}\right)^{c \lg n}}_{\text{heads}} \cdot \underbrace{\left(\frac{1}{2}\right)^{9c \lg n}}_{\text{tails}}$$

$$\text{Pr}[\text{at most } c \lg n \text{ heads}] \leq \underbrace{\binom{10c \lg n}{c \lg n}}_{\substack{\text{overestimate} \\ \text{on orders}}} \cdot \underbrace{\left(\frac{1}{2}\right)^{9c \lg n}}_{\text{tails}}$$

– Recall bounds on  $\binom{y}{x}$ :

$$\left(\frac{y}{x}\right)^x \leq \binom{y}{x} \leq \left(e \frac{y}{x}\right)^x$$

[Michael's “deathbed” formula: even on your deathbed, if someone gives you a binomial and says “simplify”, you should know this!]



- Applying this formula to the previous equation:

$$\begin{aligned}
 \Pr[\text{at most } c \lg n \text{ heads}] &\leq \binom{10c \lg n}{c \lg n} \left(\frac{1}{2}\right)^{9c \lg n} && (\beta-1)c \\
 &\leq \left(\frac{e \cdot 10c \lg n}{c \lg n}\right)^{c \lg n} \cdot \left(\frac{1}{2}\right)^{9c \lg n} && \left(\beta \cdot c \lg n\right) \\
 &= (10e)^{c \lg n} \cdot \left(\frac{1}{2}\right)^{9c \lg n} \\
 &= 2^{\lg(10e) \cdot c \lg n} \cdot \left(\frac{1}{2}\right)^{9c \lg n} \\
 &= 2^{(\lg(10e)-9)c \lg n} \\
 &= 2^{-\alpha \lg n} && \beta (\log_2(\beta e) - \beta - 1)c \\
 &= 1/n^\alpha
 \end{aligned}$$

- The point here is that, as  $10 \rightarrow \infty$ ,  $\alpha = 9 - \lg(10e) \rightarrow \infty$ , independent of (for all)  $c$

- End of proof of claim and theorem

## Acknowledgments

The mysterious “Michael” is Michael Bender at SUNY Stony Brook. This lecture is based on discussions with him.

Activity: Thought experiment — for what kind of queries/inserts would the skip list be faster than, say, an RB tree?

— finger searching: position  $d$  — can you think of a scheme to speed this up?

into A2: · implement the skip list with finger searching  
· compare against built-in ordered maps.