



L3: (a,b)-Tress

- Binary Search Tree
- AA, Left-leaning RB tree, Red-Black Tree
- B-Tree & B<sup>+</sup>-Tree

Defining (a,b) trees. (Many variants out there)

- $b \geq 2a-1$  and  $a \geq 2$
- Balanced trees
- Every non-root node has arity between a and b (inclusive)
- The root has  $\leq b$  children and at least 2 children
- All leaves are at the same depth

leaves →



Lemma: An (a,b)-tree with n nodes has height  $O(\log_a n)$ .


"Proof": In the "sparsest" tree, every inner node has a children.

$$1 + a + a^2 + a^3 + \dots + a^h \approx n \Rightarrow 1 + 2 \frac{a^{h+1} - 1}{a - 1} = n$$

$$\Rightarrow h = O(\log_a n).$$




Insertion: To insert key k: ① search for k (will end up at a leaf & fall out)  
② insert k into this leaf ③ if overflowing, restructure!



too full, split into two nodes  
#keys = b  $\Rightarrow$  promote 1 & have  $\lceil \frac{b-1}{2} \rceil$  &  
promote the middle  $\left( \left\lfloor \frac{b-1}{2} \right\rfloor \text{ left} \right)$   
 $\frac{b-1}{2} \geq \frac{2a}{2} = a$

Ex:



\* possible height increase!

Deletion. ① Locate the key we wish to delete

② If not a leaf, find the replacement node (successor)

③ Recursively delete the replacement node

°. WLOG: restructuring will begin at a leaf.

[Two cases & strategies: repair because after deletion, node is too sparse.

① Steal from sibling

if  $\exists$  a sibling to steal from.


$$d_{me} = a - 1 + 1 = a \quad \checkmark$$

stolen

② Merge with sibling.


if sibling is too sparse to steal from

$$\alpha' = a \quad (\text{si'b}) \quad + \quad \alpha = a - 1 \quad (\text{me}) \Rightarrow d' + d = 2a - 1 \leq b. \quad \checkmark$$



\* possible height decrease.


Ex:



Remarks: • (2,3) or (2,4) trees aren't binary

• What's the deal?  $K[]$  keys &  $\text{Node}[]$  children,  $\Rightarrow$  wasted space  
2 or 3 3 or 4 . weird logic

• What if they can be "seen" & stored as BSTs?



Red-Black tree (1978)  
 • Guibas & Sedgwick



• AA tree (after its inventor Arne Anderson, 1993)

$\Rightarrow$  if a node is red, it's the right child of a black node (rule out (b) above)

$\Rightarrow$  (2,3)-tree

• LLRB (Sedgwick, 2008)

Left leaning Red-Black tree



Discussion: ①  $(a, b)$  is super configurable

② Low values "reconstruct" BBT

③ Any good use for high values of  $a$  &  $b$ ?

Detour


Architecture 101:

- cache line size (32, 64, 128 bytes)  $\leftrightarrow$  size

- storage block size

eg. SSD: 512KB?

Older SATA/IDE: 4KB?



time  $\sim$  # of blocks/cache lines "touched."

B-Tree: • (B)BST  $\rightarrow$  too deep (i.e. touching too many blocks)

- Each node doesn't fully utilize the width of a line/block.
- Invented by Bayer and McCreight (Boeing Research) in 1970.
- $(a, b)$  tree where  $a = \lceil b/2 \rceil$  (i.e.,  $b = 2a$  or  $b = 2a - 1$ )
- Typical  $b \approx 100$

**Theorem 1** Starting with an empty tree, the amortized number of rebalancing operations in an  $(a, b)$ -tree is:

1.  $\Theta(\log_a n)$  when  $b = 2a - 1$ .
2.  $\Theta(1)$  when  $b = 2a$ .
3.  $\Theta(1/(\epsilon a))$  when  $b = (2 + \epsilon)a$ , for  $\epsilon > 0$ .

Applications

- File systems  $\begin{cases} \text{NTFS} \\ \text{ext4} \\ \text{APFS} \end{cases}$
- DB indexes
- NVRAM layout

B<sup>+</sup>-Tree: idea • Keep data only in the leaves  
• Chaining of the leaves.  
Why?

\* Because B+ trees don't have data associated with interior nodes, more keys can fit on a page of memory. Therefore, it will require fewer cache misses in order to access data that is on a leaf node.

\* The leaf nodes of B+ trees are linked, so doing a full scan of all objects in a tree requires just one linear pass through all the leaf nodes. A B tree, on the other hand, would require a traversal of every level in the tree. This full-tree traversal will likely involve more cache misses than the linear traversal of B+ leaves.